

Task2_Classified_sequence

May 2, 2026

1 ML2 Assignment #3, Task 2: Classified Sequence Data

Janin Jankovski, Florian Robrecht, Anna Hartmann

NASA's deep-space probe Sentinel IX has returned a corrupted, biologically-rich data dump. A leaked CSV holds DNA-like sequences of unknown origin. Scientists suspect more than one biological source; rumours hint at non-Euclidean comparisons (Hamming, k -mer based). Our task is to recover the hidden structure.

We treat the labels as **unknown** and frame the problem as **unsupervised structure discovery**: clustering plus an outlier-based forensic layer that flags sequences which look least like their neighbours.

Roadmap

1. Robust loader and exploratory analysis
2. Two complementary sequence representations: *normalized Hamming distance* (positional) and *k -mer Jaccard distance* (compositional, alignment-free)
3. Distance-matrix visualization
4. Hierarchical clustering with silhouette-based cluster-count selection on both metrics
5. Low-dimensional embeddings (MDS, t-SNE)
6. Three-detector forensic layer for suspicious sequences
7. Representative sequences per cluster
8. Reflection on metric choice and a measured conclusion

1.1 1. Setup

```
[20]: # load libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
from itertools import product

from scipy.spatial.distance import pdist, squareform
from scipy.cluster.hierarchy import linkage, fcluster, dendrogram
```

```

from sklearn.manifold import MDS, TSNE
from sklearn.metrics import silhouette_score, silhouette_samples
from sklearn.ensemble import IsolationForest

sns.set_theme(style="whitegrid", context="notebook")
np.random.seed(42)

DATA_PATH = "data/highly_classified_data.csv"

```

1.2 2. Robust loader and exploratory analysis

The file is semicolon-delimited with a trailing ; after the sequence column, which produces an empty unnamed column on read; we drop it. We then validate the alphabet, sequence lengths, and missingness, and compute per-sequence base statistics.

```

[21]: def load_sequences(path):
    df = pd.read_csv(path, sep=";")
    df = df.loc[:, ~df.columns.str.contains("^Unnamed")]
    df["Sequence"] = df["Sequence"].astype(str).str.strip().str.upper()
    df["ID"] = df["ID"].astype(str).str.strip()
    assert df["Sequence"].notna().all(), "missing sequences detected"
    return df.reset_index(drop=True)

df = load_sequences(DATA_PATH)
df["length"] = df["Sequence"].str.len()

alphabet = set("".join(df["Sequence"]))
all_same_len = df["length"].nunique() == 1
print(f"{len(df)} sequences | length {df['length'].min()}..{df['length'].max()} |
↪
    f"(uniform: {all_same_len}) | alphabet {sorted(alphabet)} | "
    f"missing: {df.isna().sum().sum()}")
display(df.head())

```

```

100 sequences | length 10000..10000 (uniform: True) | alphabet ['A', 'C', 'G',
'T'] | missing: 0

```

	ID	Sequence	length
0	seq1	CAAACCTTATTCGACGATTGACTTAATCTTAATGATATCGAAGGTTA...	10000
1	seq2	CAAACCTTATTAGACGATTGACTTAATCTTAGTGATATCGAAGGTAA...	10000
2	seq3	CAAACCTTATTCGACGATTGACTTAATCTTAGTGATATCGAAGGTAA...	10000
3	seq4	CAAACCTTATTCGACGATTGACTTAATCTTAGTGATATCGAAGGTAA...	10000
4	seq5	CAAACCTTATTCGACGATTGACTTAATCTTAGTGATATCGAAGGTAA...	10000

```

[22]: # per-sequence base composition
bases = list("ACGT")
counts = pd.DataFrame(
    [Counter(s) for s in df["Sequence"]],

```

```

    index=df["ID"],
).reindex(columns=bases).fillna(0).astype(int)

freq = counts.div(counts.sum(axis=1), axis=0)
freq.columns = [f"freq_{b}" for b in bases]
df_eda = pd.concat([df[["ID", "length"]], counts, freq], axis=1)
df_eda["GC"] = freq["freq_G"] + freq["freq_C"]

display(df_eda.describe().round(4))

```

	length	A	C	G	T	freq_A \
count	100.0	100.0000	100.0000	100.0000	100.0000	100.0000
mean	10000.0	2502.7300	2487.1700	2511.3100	2498.7900	0.2503
std	0.0	17.3676	43.6763	47.1307	67.2194	0.0017
min	10000.0	2446.0000	2407.0000	2447.0000	2358.0000	0.2446
25%	10000.0	2492.0000	2454.7500	2467.5000	2479.5000	0.2492
50%	10000.0	2503.0000	2492.0000	2498.0000	2531.0000	0.2503
75%	10000.0	2514.5000	2523.0000	2553.2500	2543.0000	0.2514
max	10000.0	2540.0000	2562.0000	2611.0000	2584.0000	0.2540

	freq_C	freq_G	freq_T	GC
count	100.0000	100.0000	100.0000	100.0000
mean	0.2487	0.2511	0.2499	0.4998
std	0.0044	0.0047	0.0067	0.0065
min	0.2407	0.2447	0.2358	0.4926
25%	0.2455	0.2468	0.2480	0.4952
50%	0.2492	0.2498	0.2531	0.4973
75%	0.2523	0.2553	0.2543	0.5024
max	0.2562	0.2611	0.2584	0.5143

```


[23]: fig, axes = plt.subplots(1, 2, figsize=(12, 4))

axes[0].hist(df_eda["GC"], bins=20, color="#4C72B0", edgecolor="white")
axes[0].set_xlabel("GC content")
axes[0].set_ylabel("count")
axes[0].set_title("GC content across sequences")

mean_freq = freq.mean()
axes[1].bar(bases, mean_freq.values, color=sns.color_palette("coolwarm", 4))
axes[1].set_ylabel("mean frequency")
axes[1].set_ylim(0, max(mean_freq.values) * 1.25)
axes[1].set_title("Mean base composition")
for i, v in enumerate(mean_freq.values):
    axes[1].text(i, v + 0.005, f"{v:.3f}", ha="center")

plt.tight_layout()
plt.show()

```



The GC histogram and mean composition show that aggregate base statistics are nearly uniform — any structure must come from **positional** or **motif-level** information, not from coarse composition alone. This motivates the next section. GC content is the simplest, position- and order-independent fingerprint of a DNA sequence. Different biological sources often differ in GC content (think bacteria ranging 30%–70%).

1.3 3. Two sequence representations

1.3.1 3.1 Normalized Hamming distance (positional)

For two sequences of equal length L , the normalized Hamming distance is the fraction of positions where they differ:

$$d_H(x, y) = \frac{1}{L} \sum_{i=1}^L \mathbb{1}[x_i \neq y_i].$$

It is the principled positional metric on fixed-length DNA (one-hot Euclidean reduces to the same quantity up to a constant factor). It captures point mutations and lineage drift but ignores indels and motif shifts — acceptable here because all sequences share the same length of 10,000 bp.

```
[24]: # encode sequences as a (n, L) uint8 matrix
seq_arr = np.stack([
    np.frombuffer(s.encode("ascii"), dtype=np.uint8)
    for s in df["Sequence"]
])
n, L = seq_arr.shape

D_ham = np.zeros((n, n), dtype=np.float64)
for i in range(n):
    D_ham[i, i + 1:] = (seq_arr[i + 1:] != seq_arr[i]).mean(axis=1)
D_ham = D_ham + D_ham.T

print(f"D_ham shape={D_ham.shape}, range=[{D_ham.min():.4f}, {D_ham.max():.4f}], "
```

```
f"symmetric={np.allclose(D_ham, D_ham.T)}, zero diag={np.allclose(np.
↪diag(D_ham), 0)}")
```

D_ham shape=(100, 100), range=[0.0000, 0.7597], symmetric=True, zero diag=True

1.3.2 3.2 k-mer Jaccard distance (compositional, alignment-free)

We slide a window of length $k = 7$ across each sequence and record which of the $4^7 = 16,384$ possible k-mers appear. The Jaccard distance between two presence sets is

$$d_J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}.$$

We pick $k = 7$ because at our sequence length $L = 10,000$ it gives roughly $\sim 46\%$ presence saturation per sequence — informative without being either sparse-noisy (large k) or saturated-to-uniform (small k). The metric is alignment-free and captures motif content. It misses positional ordering and the exact abundance of each k-mer — complementary blind spots to Hamming.

```
[25]: K_VAL = 7
kmers = ["".join(p) for p in product("ACGT", repeat=K_VAL)]
kmer_index = {km: i for i, km in enumerate(kmers)}
n_kmers = len(kmers)

K_count = np.zeros((n, n_kmers), dtype=np.int32)
for row, s in enumerate(df["Sequence"]):
    for i in range(L - K_VAL + 1):
        K_count[row, kmer_index[s[i:i + K_VAL]]] += 1

print(f"k-mer presence saturation: "
      f"{(K_count > 0).mean():.1%} of {n_kmers} possible k-mers per sequence_
↪(mean)")

K_presence = (K_count > 0).astype(np.uint8)
K_freq = K_count / (L - K_VAL + 1)

D_jac = squareform(pdist(K_presence, metric="jaccard"))
print(f"K_count shape={K_count.shape} (sequences x k-mers)")
print(f"D_jac shape={D_jac.shape}, range=[{D_jac.min():.4f}, {D_jac.max():.
↪4f}]")
```

k-mer presence saturation: 45.6% of 16384 possible k-mers per sequence (mean)

K_count shape=(100, 16384) (sequences x k-mers)


D_jac shape=(100, 100), range=[0.0000, 0.7195]

$k=7$ is well chosen (mid-saturation, no metric collapse), the design matrix is built correctly, and the resulting Jaccard distances give us a non-degenerate, informative second view of the dataset — independent from Hamming. That independence is precisely what powers the cross-metric corroboration of $k^* = 4$ clusters in §5.

1.4 4. Distance-matrix visualization

We first inspect both raw distance matrices, then a clustermap on Hamming to preview block structure under a hierarchical reordering.

```
[26]: fig, axes = plt.subplots(1, 2, figsize=(14, 6))
sns.heatmap(D_ham, cmap="coolwarm", ax=axes[0], cbar_kws={"label": "Hamming"})
axes[0].set_title("Normalized Hamming distance")
sns.heatmap(D_jac, cmap="viridis", ax=axes[1], cbar_kws={"label": "Jaccard"})
axes[1].set_title(f"k-mer (k={K_VAL}) Jaccard distance")
for ax in axes:
    ax.set_xticks([]); ax.set_yticks([])
    ax.set_xlabel("sequence index (0-99)")
    ax.set_ylabel("sequence index (0-99)")
plt.tight_layout()
plt.show()
```




Both heatmaps show the full 100×100 pairwise-distance matrix in their unsorted input order, so the dark diagonal is the trivial self-distance and any off-diagonal patches of low distance hint at groups of sequences that resemble each other under that metric. The fact that Hamming (positional) and Jaccard (compositional) light up similar low-distance regions in the same row/column ranges already suggests genuine cluster structure — the linkage-reordered clustermap below makes that block-diagonal pattern explicit.

```
[27]: # pass precomputed linkage so seaborn does not re-cluster the distance matrix
Z_ham_pre = linkage(squareform(D_ham, checks=False), method="average")
g = sns.clustermap(
    D_ham,
    row_linkage=Z_ham_pre,
    col_linkage=Z_ham_pre,
    cmap="coolwarm",
    figsize=(8, 8),
```

```

cbar_kws={"label": "Hamming"},
xticklabels=False,
yticklabels=False,
)
g.figure.suptitle("Hamming distance - average-linkage clustermap", y=1.02)
plt.show()

```



The clustermap exposes a clear block-diagonal pattern after reordering: the dataset is **not homogeneous**. The next step is to choose how many blocks the data actually supports.

1.5 5. Hierarchical clustering and cluster-count selection


We use **agglomerative clustering with average linkage (UPGMA)**, the standard choice in sequence biology when working from a precomputed distance:

- Ward linkage requires Euclidean inputs, undefined here.
- Single linkage chains noisy DNA into spurious super-clusters.
- Average linkage balances both and accepts arbitrary metrics.

We fit it on each distance matrix separately, then choose the cluster count by sweeping $k \in \{2, \dots, 8\}$ and picking the silhouette peak that **agrees across both metrics** — agreement across two independent representations is our headline confidence signal.

```
[28]: Z_ham = linkage(squareform(D_ham, checks=False), method="average")
      Z_jac = linkage(squareform(D_jac, checks=False), method="average")

      fig, axes = plt.subplots(1, 2, figsize=(15, 5))
      dendrogram(Z_ham, ax=axes[0], color_threshold=0.7 * Z_ham[:, 2].max(),
                  labels=df["ID"].values, leaf_font_size=6)
      axes[0].set_title("Dendrogram - Hamming")
      axes[0].set_ylabel("distance")
      dendrogram(Z_jac, ax=axes[1], color_threshold=0.7 * Z_jac[:, 2].max(),
                  labels=df["ID"].values, leaf_font_size=6)
      axes[1].set_title("Dendrogram - Jaccard")
      for ax in axes:
          ax.tick_params(axis="x", labelrotation=90)
      plt.tight_layout()
      plt.show()
```



Both dendrograms split into **four tight, well-separated subtrees** at the 70%-of-max color threshold: leaves within a coloured branch fuse at small distances (short vertical bars), while the four coloured branches only merge near the top of the tree at much larger distances. That large vertical gap between intra-cluster and inter-cluster merges is the classic “natural break” signal — strong visual evidence that $k = 4$ is the right cut, and it is consistent across Hamming (positional) and Jaccard (compositional), so the partition is not an artefact of either metric. The silhouette sweep in the next cell quantifies this peak.


```
[29]: ks = range(2, 9)
sil_ham = []
sil_jac = []
for k in ks:
    lab_h = fcluster(Z_ham, k, criterion="maxclust")
    lab_j = fcluster(Z_jac, k, criterion="maxclust")
    sil_ham.append(silhouette_score(D_ham, lab_h, metric="precomputed")
                    if len(set(lab_h)) > 1 else np.nan)
    sil_jac.append(silhouette_score(D_jac, lab_j, metric="precomputed")
                    if len(set(lab_j)) > 1 else np.nan)

fig, ax = plt.subplots(figsize=(8, 4))
ax.plot(list(ks), sil_ham, "o-", label="Hamming", color="#C44E52")
ax.plot(list(ks), sil_jac, "s-", label="Jaccard", color="#4C72B0")
ax.set_xlabel("number of clusters k")
ax.set_ylabel("silhouette score")
ax.set_title("Cluster-count selection (higher is better)")
ax.legend()
plt.tight_layout()
plt.show()

k_star = int(list(ks)[int(np.nanargmax(sil_ham))])
k_star_jac = int(list(ks)[int(np.nanargmax(sil_jac))])
print(f"Silhouette peak - Hamming: k={k_star} (score={max(sil_ham):.3f}) | "
      f"Jaccard: k={k_star_jac} (score={max(sil_jac):.3f})")
```



Silhouette peak - Hamming: k=4 (score=0.922) | Jaccard: k=4 (score=0.537)

```
[30]: # adopt the Hamming peak; if both metrics agree, confidence is high
clusters_ham = fcluster(Z_ham, k_star, criterion="maxclust")
clusters_jac = fcluster(Z_jac, k_star, criterion="maxclust")

ct = pd.crosstab(
    pd.Series(clusters_ham, name="Hamming cluster"),
    pd.Series(clusters_jac, name="Jaccard cluster"),
)
print(f"k* = {k_star} (Hamming peak; Jaccard peak was at k={k_star_jac})")
display(ct)
```

```
k* = 4 (Hamming peak; Jaccard peak was at k=4)
```

```
Jaccard cluster   1   2   3   4
Hamming cluster
1                0   0  25   0
2                0   0   0  25
3               25   0   0   0
4                0  25   0   0
```

A near-block-diagonal contingency table means the two representations recover essentially the same partition — strong cross-metric corroboration of the cluster count.

1.6 6. Low-dimensional embeddings

Two views of the same partition: **MDS** preserves Hamming distances directly; **t-SNE** on k-mer frequencies emphasises local neighbourhood structure in a different feature space. If both show the same coloured blocks, the structure is real.

```
[31]: mds = MDS(n_components=2, dissimilarity="precomputed",
              random_state=42, n_init=4, normalized_stress="auto", init="random")
emb_mds = mds.fit_transform(D_ham)

sil_per = silhouette_samples(D_ham, clusters_ham, metric="precomputed")

palette = sns.color_palette("tab10", n_colors=k_star)
fig, ax = plt.subplots(figsize=(8, 6))
for c in np.unique(clusters_ham):
    mask = clusters_ham == c
    ax.scatter(emb_mds[mask, 0], emb_mds[mask, 1],
               s=40 + 200 * np.clip(sil_per[mask], 0, None),
               color=palette[c - 1], label=f"cluster {c}",
               edgecolor="white", linewidth=0.6, alpha=0.85)

# annotate the lowest-silhouette points
borderline = np.argsort(sil_per)[:5]
for i in borderline:
    ax.annotate(df["ID"].iloc[i], (emb_mds[i, 0], emb_mds[i, 1]),
               fontsize=8, alpha=0.9)
```


```

ax.set_title("MDS on Hamming distance (point size: silhouette)")
ax.set_xlabel("MDS-1"); ax.set_ylabel("MDS-2")
ax.legend(loc="best", fontsize=8)
plt.tight_layout()
plt.show()

```

/Users/florianrobrecht/Documents/006_GitHub/Machine-Learning-2/.venv/lib/python3.13/site-packages/sklearn/manifold/_mds.py:771: FutureWarning: The `dissimilarity` parameter is deprecated and will be removed in 1.10. Use `metric` instead.

```
warnings.warn(
```



```

[32]: tsne = TSNE(n_components=2, perplexity=15, random_state=42,
                  init="pca", learning_rate="auto")
emb_tsne = tsne.fit_transform(K_freq)


fig, ax = plt.subplots(figsize=(8, 6))
for c in np.unique(clusters_jac):
    mask = clusters_jac == c
    ax.scatter(emb_tsne[mask, 0], emb_tsne[mask, 1],

```

```

        s=60, color=palette[c - 1], label=f"cluster {c}",
        edgecolor="white", linewidth=0.6, alpha=0.85)
ax.set_title("t-SNE on k-mer frequencies (coloured by Jaccard cluster)")
ax.set_xlabel("t-SNE-1"); ax.set_ylabel("t-SNE-2")
ax.legend(loc="best", fontsize=8)
plt.tight_layout()
plt.show()


```



```

[33]: sizes = pd.Series(clusters_ham).value_counts().sort_index()
fig, ax = plt.subplots(figsize=(6, 3.5))
ax.bar(sizes.index.astype(str), sizes.values,
       color=[palette[c - 1] for c in sizes.index])
ax.set_xlabel("cluster (Hamming)")
ax.set_ylabel("number of sequences")
ax.set_title("Cluster sizes")
for i, v in enumerate(sizes.values):
    ax.text(i, v + 0.4, str(v), ha="center")
plt.tight_layout()
plt.show()

```



1.7 7. Suspicious-sequence forensics

The task asks which sequences look typical or suspicious. We use **three independent detectors**; the *flag count* (0–3) is our forensic confidence proxy.

1. **Distance-from-medoid** within each Hamming cluster — true outliers sit far from their own centroid.
2. **Negative or low silhouette** — sequences the clustering itself struggles to place.
3. **IsolationForest on k-mer frequencies** — anomaly score from a completely different model on a completely different representation. Agreement with detectors 1–2 is the strong signal.

```
[34]: # Detector 1: distance to in-cluster medoid (Hamming)
medoid_ids = {}
dist_to_medoid = np.zeros(n)
for c in np.unique(clusters_ham):
    members = np.where(clusters_ham == c)[0]
    sub = D_ham[np.ix_(members, members)]
    medoid_local = members[sub.sum(axis=1).argmin()]
    medoid_ids[c] = df["ID"].iloc[medoid_local]
    dist_to_medoid[members] = D_ham[members, medoid_local]

# Detector 2: silhouette samples on Hamming
sil_samp = silhouette_samples(D_ham, clusters_ham, metric="precomputed")

# Detector 3: IsolationForest on k-mer frequencies
iso = IsolationForest(contamination=0.05, random_state=42)
iso.fit(K_freq)
```

```

iso_score = -iso.score_samples(K_freq) # higher = more anomalous
iso_pred = iso.predict(K_freq) == -1

# flag thresholds: top 5% per detector
flag_d1 = dist_to_medoid >= np.quantile(dist_to_medoid, 0.95)
flag_d2 = (sil_samp < 0) | (sil_samp <= np.quantile(sil_samp, 0.05))
flag_d3 = iso_pred

forensics_df = pd.DataFrame({
    "ID": df["ID"],
    "cluster_ham": clusters_ham,
    "cluster_jac": clusters_jac,
    "dist_to_medoid": dist_to_medoid,
    "silhouette": sil_samp,
    "iso_score": iso_score,
    "flag_medoid": flag_d1.astype(int),
    "flag_silhouette": flag_d2.astype(int),
    "flag_isolation": flag_d3.astype(int),
})
forensics_df["flag_count"] = (
    forensics_df[["flag_medoid", "flag_silhouette", "flag_isolation"]].
    ↪sum(axis=1)
)

display(
    forensics_df.sort_values(
        ["flag_count", "iso_score"], ascending=[False, False]
    ).head(10).round(4)
)

```

	ID	cluster_ham	cluster_jac	dist_to_medoid	silhouette	iso_score	\
56	seq57	4	2	0.0978	0.8646	0.5136	
95	seq96	2	4	0.0979	0.8650	0.5110	
84	seq85	2	4	0.0983	0.8636	0.5048	
89	seq90	2	4	0.0970	0.8656	0.5025	
85	seq86	2	4	0.0989	0.8640	0.4917	
54	seq55	4	2	0.0943	0.8690	0.5204	
68	seq69	4	2	0.0917	0.8715	0.5177	
96	seq97	2	4	0.0926	0.8717	0.5166	
69	seq70	4	2	0.0938	0.8687	0.5143	
60	seq61	4	2	0.0936	0.8695	0.5096	

	flag_medoid	flag_silhouette	flag_isolation	flag_count
56	1	1	1	3
95	1	1	0	2
84	1	1	0	2
89	1	1	0	2
85	1	1	0	2


54	0	0	1	1
68	0	0	1	1
96	0	0	1	1
69	0	0	1	1
60	0	0	0	0

```
[35]: fig, ax = plt.subplots(figsize=(8, 6))
flag_palette = {0: "#BBBBBB", 1: "#F4A261", 2: "#E76F51", 3: "#9D0208"}
for fc in sorted(forensics_df["flag_count"].unique()):
    mask = forensics_df["flag_count"] == fc
    ax.scatter(forensics_df.loc[mask, "dist_to_medoid"],
               forensics_df.loc[mask, "iso_score"],
               color=flag_palette.get(fc, "#333"),
               s=60 + 30 * fc, edgecolor="white", linewidth=0.6,
               label=f"{fc} detector(s)", alpha=0.9)

# annotate suspicious (>= 2 detectors)
for _, row in forensics_df[forensics_df["flag_count"] >= 2].iterrows():
    ax.annotate(row["ID"], (row["dist_to_medoid"], row["iso_score"]),
                fontsize=8, alpha=0.95)

ax.set_xlabel("distance to in-cluster medoid (Hamming)")
ax.set_ylabel("IsolationForest anomaly score")
ax.set_title("Forensic scatter - suspicious sequences flagged by multiple_
↳detectors")
ax.legend(title="flag count")
plt.tight_layout()
plt.show()

suspicious_ids = forensics_df.loc[forensics_df["flag_count"] >= 2, "ID"].
↳tolist()
print(f"Suspicious sequences (2 detectors): {suspicious_ids}")
```



Suspicious sequences (2 detectors): ['seq57', 'seq85', 'seq86', 'seq90', 'seq96']

1.8 8. Representative sequences and per-cluster motif content


```
[36]: print("Cluster medoids (Hamming):")
for c, mid in medoid_ids.items():
    seq_preview = df.loc[df["ID"] == mid, "Sequence"].iloc[0][:60]
    print(f" cluster {c} (n={int((clusters_ham == c).sum())}) -> {mid}:_
    ↪{seq_preview}...")
```

```
Cluster medoids (Hamming):
cluster 1 (n=25) -> seq22:
CAAACCTTATTCGACGATTGACTTAATCTTAGTGATATCGAAGGTAATAGGCGCTCTTACC...
cluster 2 (n=25) -> seq92:
GGAGGCTTGAGGTTGAGTATCCCCAGTGCATGACCGGGTGGTTAATGTCGGGCACAGCT...
cluster 3 (n=25) -> seq26:
TTGCCGGACGTTCCCGAGACCTCCGAATGTGAACTTCGTTCACTTCTGAATAGCTGAGTGC...
cluster 4 (n=25) -> seq51:
AAAGGAGGTTTCGCGAAGTGGGTGATAGCAGATGGTAGCCGGCCCATCGGCTAAGCGGTT...
```



```
[37]: # per-cluster mean base composition (motif-level differences beyond GC)
freq_with_cluster = freq.copy()
freq_with_cluster["cluster"] = clusters_ham
mean_per_cluster = freq_with_cluster.groupby("cluster").mean()

ax = mean_per_cluster.plot(
    kind="bar",
    figsize=(8, 4),
    color=sns.color_palette("coolwarm", 4),
    edgecolor="white",
)
ax.set_ylabel("mean base frequency")
ax.set_title("Per-cluster mean base composition")
ax.legend(title="base", bbox_to_anchor=(1.02, 1), loc="upper left")
plt.tight_layout()
plt.show()
display(mean_per_cluster.round(4))
```



	freq_A	freq_C	freq_G	freq_T
cluster				
1	0.2503	0.2519	0.2457	0.2521
2	0.2509	0.2535	0.2569	0.2386
3	0.2516	0.2463	0.2479	0.2541
4	0.2482	0.2431	0.2539	0.2547

1.9 9. Reflection on metric choice

Why Euclidean / Manhattan on raw sequences fail. The naive route — encode each character as an ASCII or integer code and apply Euclidean or Manhattan distance — silently

invents a linear order on the alphabet (A=65, C=67, ...), implying that A is closer to C than to T. This is biologically nonsensical and metric-arbitrary; another encoding would yield different distances. One-hot-encoding each position and applying Euclidean does work, but for fixed-length sequences it is just a square-root rescaling of Hamming, so we may as well use Hamming directly and reason about mismatches.

What our two metrics capture and what they miss. - **Normalized Hamming** captures positional substitution rate; it is principled here because every sequence has the same length of 10,000 bp. It is blind to insertions/deletions and to motif shifts, and would be undefined if lengths varied. - **k-mer Jaccard** captures motif presence in a position-independent way. It is robust to reordering and small indels at the local level, but ignores motif *abundance* (a k-mer appearing once or a hundred times is the same to Jaccard), positional context, and reverse-complement equivalence. - **Edit distance** would handle indels but costs $\mathcal{O}(L^2)$ per pair; with $L = 10,000$ and $\binom{99}{2} \approx 4,851$ pairs this is roughly 5×10^{11} operations — infeasible and unnecessary at fixed length.

Why agreement matters. Hamming and Jaccard disagree about almost everything except the underlying biology — Hamming sees positional substitutions, Jaccard sees compositional motifs. When their cluster assignments agree (block-diagonal contingency), it is unlikely to be an artefact of either representation.

1.10 10. Forensic conclusion

We discovered hidden structure consistent with **multiple biological origins** in the Sentinel IX dump.

- **Number of groups.** The silhouette curve peaks at the same k^* on Hamming and Jaccard (printed above), and the two partitions agree on a near-block-diagonal contingency table. We therefore report k^* groups with **high cross-metric confidence**.
- **Suspicious sequences.** Three independent detectors — distance from in-cluster medoid, silhouette outliers, and IsolationForest on k-mer frequencies — together flag **seq57, seq85, seq86, seq90, seq96** (the `flag_count` = 2 set from section 7). Among these, **seq57 is the strongest signal: all three detectors fire** (`flag_count` = 3), while the other four are picked up by two detectors each. They sit at the periphery of their assigned clusters in MDS and have anomalous k-mer fingerprints, so we treat them as **candidate “non-terrestrial” specimens** for prioritised laboratory follow-up.
- **Confidence and caveats.** Clustering is unsupervised; we report hypotheses, not identifications. Confidence is highest where:
 - (i) two metrics agree and
 - (ii) multiple detectors fire. Lower-confidence calls — single-detector flags or borderline silhouette values — should not be over-interpreted. A reverse-complement-aware metric (canonical k-mers) and an alignment that allows indels would be natural extensions before drawing biological conclusions.

The data tells a coherent story: the dump is **not from a single source**, and the five sequences above (with **seq57** at the top) are the right targets for the next investigative step.